



DBMS

Database Management System with MySQL

Yash Lathiya

Table of Contents

Sr No.	Topic	Date	Page No.
1	Overview of DBMS		
2	MYSQL		
3	Overview of Workbench		
4	Database Design		
5	SQL Basics		
6	Data Sorting		
7	Null Value & Keyword		
8	Auto Increment		
9	Data Languages		
10	Limit		
11	Aggregate Functions		
12	Sub Queries		
13	Joins		
14	Unions		
15	Index		
16	View		
17	Backup, Restore, Explain		

Overview of DBMS

Database

A database is a collection of related data which represents some aspect of the real world. A database system is designed to be built and populated with data for a certain task.

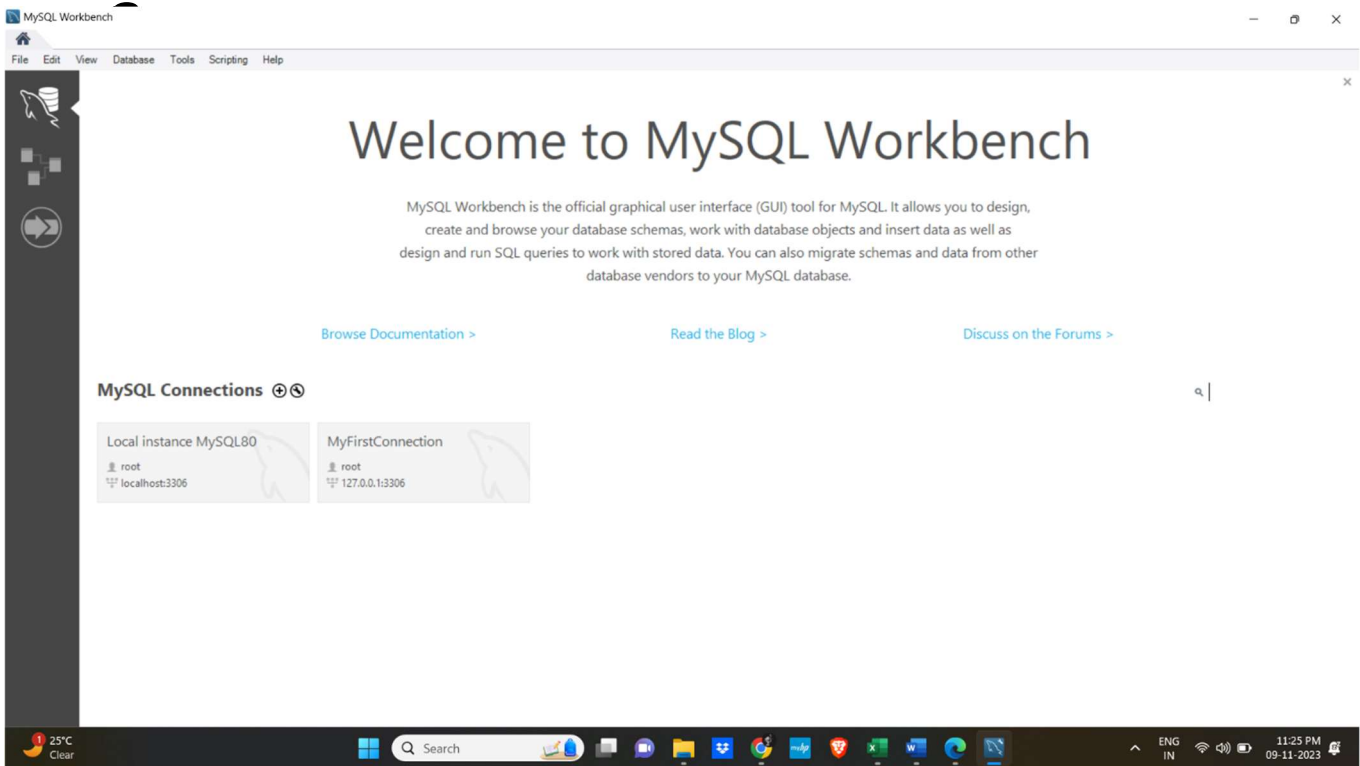
Database Management System

DBMS is a software for storing and retrieving users' data while considering appropriate security measures. It consists of a group of programs which manipulate the database. The DBMS accepts the request for data from an application and instructs the operating system to provide the specific data.

MYSQL

MySQL is structured query language which is also relational database management system. By using different types of data languages, we can insert, remove and retrieve data in efficient manner with help of MySQL.

Overview of Workbench



Workbench features

Connection

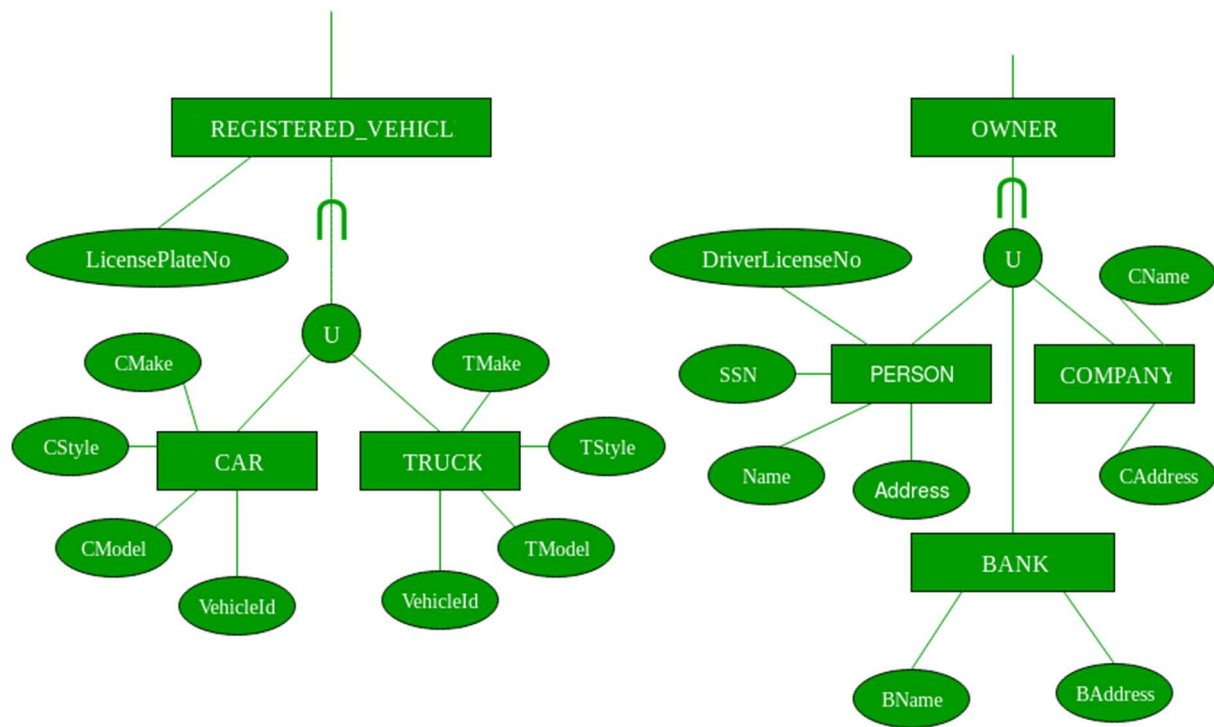
Model

Scripting

Database Design

EER Diagram

Database design can be done through applying significant constraints on database like applying primary key, foreign key, data length constraints. All these types of constraints can be visibly represented by EER diagram which is enhanced entity relation diagram.



EER diagram also includes concepts of Generalization & specialization..




Fig. Generalization




Fig. Specialization

SQL Basics

Data Sorting

-- Sorting of Data

-- Here data is sorted by ID then FirstName then LastName
in ascending order

```
SELECT
    Id,
    FirstName,
    LastName
FROM
    MyFirstTable
ORDER BY
    Id, FirstName, LastName ASC;
```

-- Here data is sorted by ID then FirstName then LastName
in descending order

```
SELECT
    Id,
    FirstName,
    LastName
FROM
    MyFirstTable
ORDER BY
    Id, FirstName, LastName DESC;
```

Null Value & Keyword

There is three type of values null values, non null values & empty values.

Null values are values which are not provided to database, It will not displayed when select query is fired.

Empty values are values which is empty string, It represents that value is not there or not provided, It will displayed when select query is fired.

Not null is keyword which implements a functionality in which specification and insertion of values are compulsory otherwise it will generate errors, It does not support null values.

By default columns are treated in such a way that it allows null values.

ALTER TABLE

MyFirstTable

ADD

Address varchar(255) NOT NULL;

Auto Increment

```
-- Create Table which contains one auto increment field  
-- Auto Increment strats from 1001 with increment of 1  
-- while both contains 1 as default values.
```

```
CREATE TABLE  
    Employee (  
        Id int PRIMARY KEY AUTO_INCREMENT,  
        FirstName varchar(255),  
        LastName varchar(255),  
        Organization varchar(255),  
        Salary int  
    ) AUTO_INCREMENT = 1001;
```


Data Languages

DDL

- Data Definition Language is a set of SQL commands used to define, modify, and manage the structure of a database.
- It includes commands like
- **CREATE** (used to create objects like tables, indexes, etc.)
- **ALTER** (used to modify existing objects)
- **DROP** (used to delete objects)
- **TRUNCATE** (used to remove all records from a table), and
- **RENAME** (used to rename objects).

DML

- Data Manipulation Language is a set of SQL commands used to manage data within a database.
- It includes commands like
- **SELECT** (used to retrieve data)
- **INSERT** (used to add new records),
- **UPDATE** (used to modify existing records),
- **DELETE** (used to remove records).

DCL

- Data Control Language is a set of SQL commands used to control access to data within a database.
- It includes commands like
- **GRANT** (used to provide specific privileges to users or roles),
- **REVOKE** (used to revoke privileges), and
- **DENY** (used to deny access).

TCL

- Transaction Control Language is a set of SQL commands used to manage transactions within a database.
- It includes commands like
- **COMMIT** (used to save changes made during the current transaction)
- **ROLLBACK** (used to undo changes made during the current transaction), and
- **SAVEPOINT** (used to set points within a transaction to which you can later roll back).

DQL

- Data Query Language is a subset of SQL that is used to query and retrieve data from a database.
- The primary command in DQL is **SELECT**, which is used to fetch data from one or more tables based on specified criteria.
- Example: **SELECT * FROM table_name WHERE condition;**

Limit

```
SELECT
    ProductID,
    ProductName,
    Price
FROM
    Products
LIMIT      -- Selecting first 50 records
    50
OFFSET     -- Neglecting first 10 records
    10;
```

Aggregate Functions

-- Count

```
SELECT
    COUNT(OrderId)
FROM
    Orders;
```

-- Sum

```
SELECT
    SUM(Price)
FROM
    Products
WHERE
    Price < 100;
```

-- Average

```
SELECT
    AVG(Price)
FROM
    Products;
```

-- Minimum

```
SELECT
    MIN(Price)
FROM
    Products
WHERE
    PRICE > 100;
```

```
-- Maximum
```

```
SELECT  
    MAX(Price)  
FROM  
    Products  
WHERE  
    PRICE < 200;
```

Sub Queries

-- Find details of product whose price is more than avg.
price


```
SELECT
    ProductID,
    ProductName,
    Price
FROM
    Products
WHERE
    ProductID
IN
    (SELECT
        ProductID
    FROM
        products
    WHERE
        price > (SELECT AVG(Price) FROM Products)
    );
```

-- Correlated Subqueries


```
SELECT
    ProductID,
    ProductName,
    Price
FROM
    Products v_product
WHERE
    Price > (SELECT
                AVG(Price)
            FROM
                Products
            );
```

There are operators like IN, NOT IN, EXIST which is used while implementing subqueries.


Joins




```
SELECT
    Orders.OrderID,
    Customers.CustomerID,
    OrderDate
FROM
    Customers
INNER JOIN
    Orders
ON
    Orders.CustomerID = Customers.CustomerID;
```

```
SELECT
    Orders.OrderID,
    Customers.CustomerID,
    OrderDate
FROM
    Customers
LEFT JOIN
    Orders
ON
    Orders.CustomerID = Customers.CustomerID;
```



```
SELECT
    Orders.OrderID,
    Customers.CustomerID,
    OrderDate
FROM
    Customers
RIGHT JOIN
    Orders
ON
    Orders.CustomerID = Customers.CustomerID;
```



```
SELECT
    Orders.OrderID,
    Customers.CustomerID,
    OrderDate
FROM
    Customers
CROSS JOIN
    Orders
ON
    Orders.CustomerID = Customers.CustomerID;
```

Unions

-- Union

It displays all unique values of dataset after performing union operation.

In simple words, It returns a set.

No duplication of values in union.

```
SELECT
    'Shipper' AS Type,
    ShipperID,
    ShipperName
FROM
    Shippers
UNION SELECT
    'Supplier' AS Type,
    SupplierID,
    SupplierName
FROM
    Suppliers;
```

-- Union All

It displays all values (with duplication) which represents data combination of both queries.

As it allows duplicate values, It may or may not be a set.

```
SELECT
    'Shipper' AS Type,
    ShipperID,
    ShipperName
FROM
    Shippers
UNION SELECT
    'Supplier' AS Type,
    SupplierID,
    SupplierName
FROM
    Suppliers;
```

Index

An index is a database object that improves the speed of data retrieval operations on a table. It serves as a mechanism to optimize query performance by allowing the database engine to quickly locate and access the rows of a table.

Purpose of Indexing:

- Faster Data Retrieval:

Indexes provide a faster way to look up and retrieve specific rows from a table, especially when dealing with large datasets.

- Query Optimization:

Indexes optimize the execution of SELECT, JOIN, and WHERE clauses in SQL queries.

CREATE INDEX

Index_ID

ON

Customers(CustomerID);

-- Unique Index

-- Doesn't allow duplicate values

CREATE UNIQUE INDEX

Index_Contact

ON

Customers(CustomerID, ContactName);

View

-- view

-- view is a virtual table which is based on sql statements and conditions.

-- view has rows and columns as real tables which can be initialized or created as well as updated and dropped

-- create view

-- France Customers

CREATE VIEW

 v_franceCustomers

AS

SELECT

 CustomerID,

 CustomerName,

 City

FROM

 Customers

WHERE

 Country = 'France';

-- How to use view

SELECT

 CustomerID,

 CustomerName,

 City

FROM

 v_franceCustomers;

```
-- UPDATE VIEW
```

```
CREATE OR REPLACE VIEW
```

```
    v_franceCustomers
```

```
AS
```

```
SELECT
```

```
    CustomerID,
```

```
    CustomerName,
```

```
    City,
```

```
    Address
```

```
FROM
```

```
    Customers
```

```
WHERE
```

```
    Country = 'France';
```

```
-- Drop(Discard) View
```

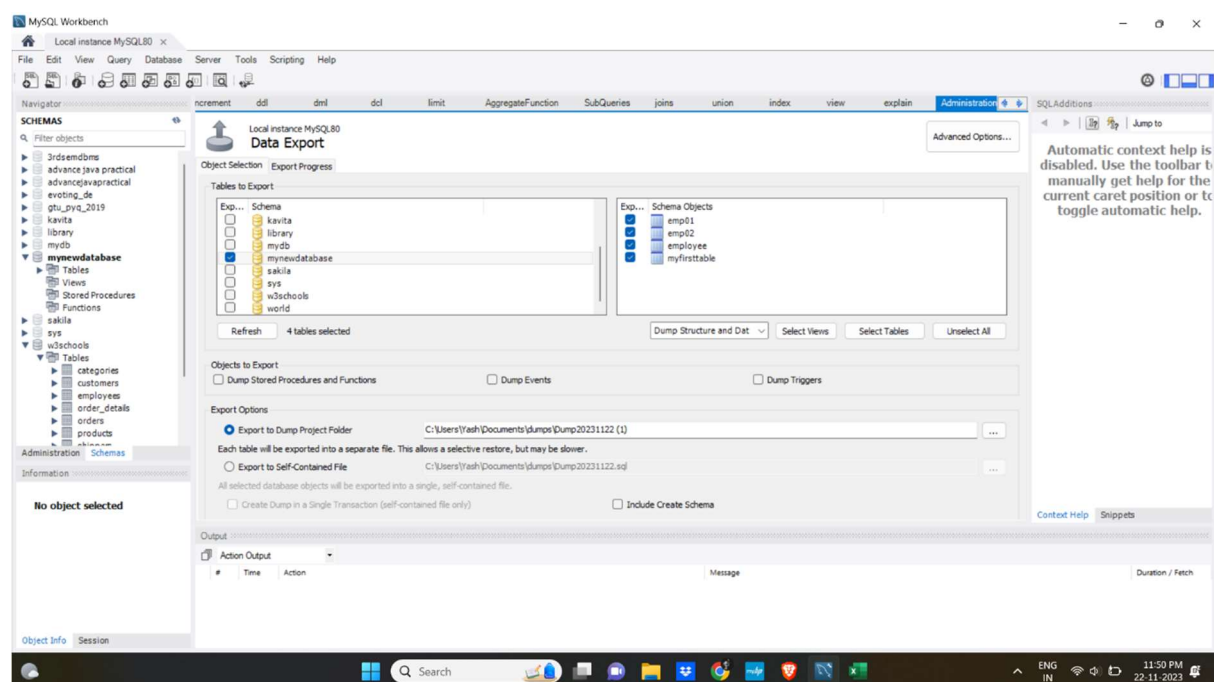
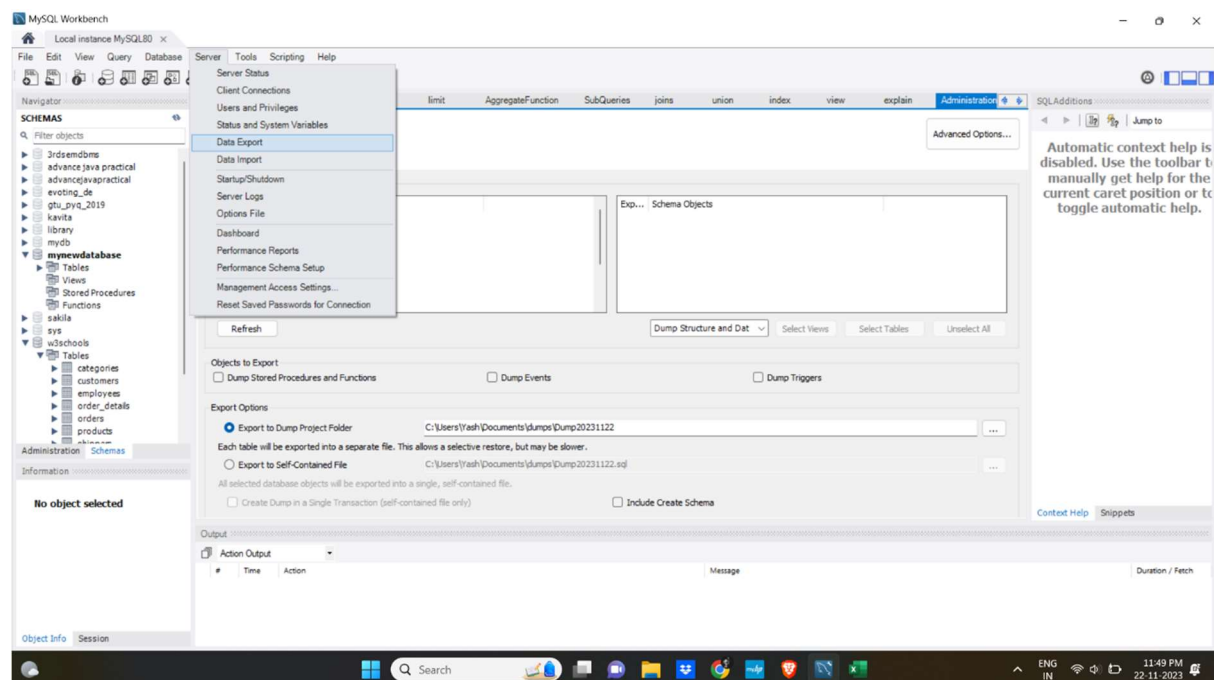
```
DROP VIEW
```

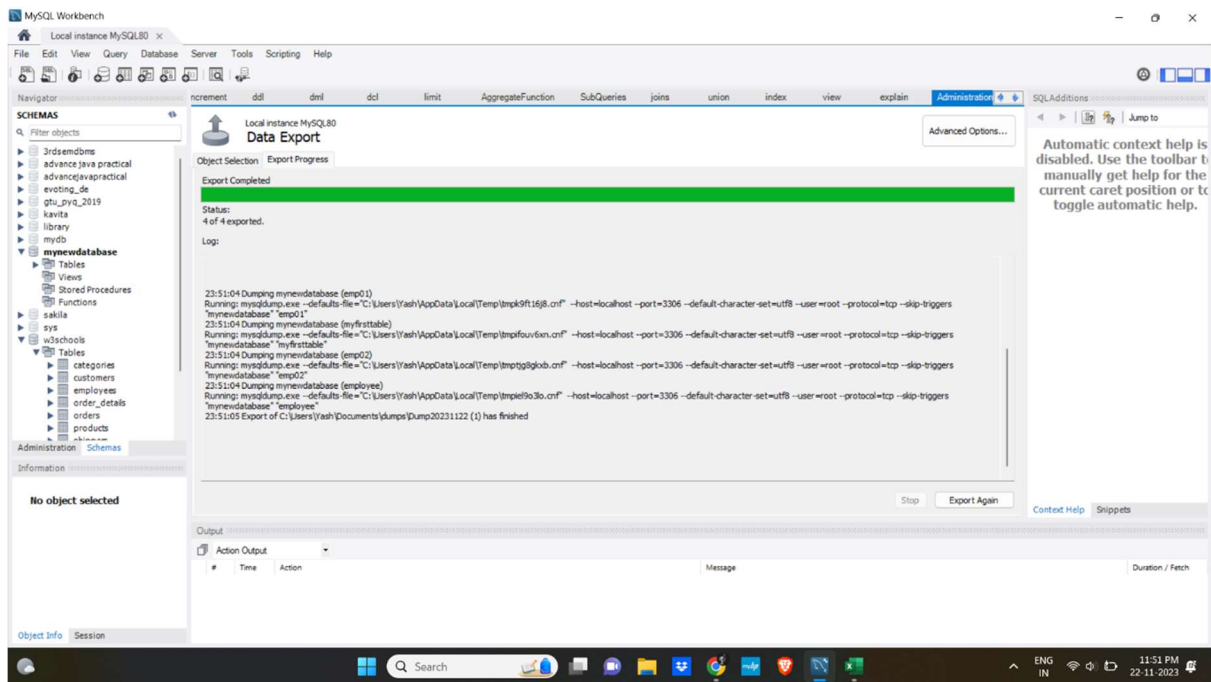
```
    v_franceCustomers;
```


Backup, Restore, Explain

Backup

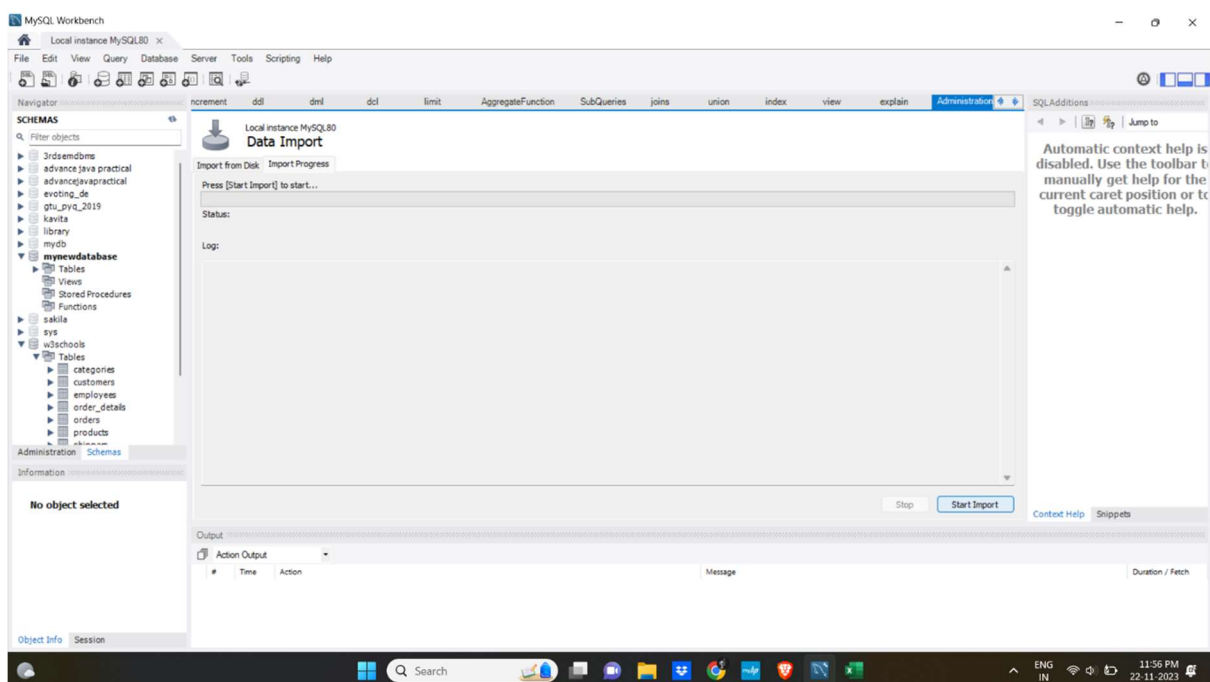
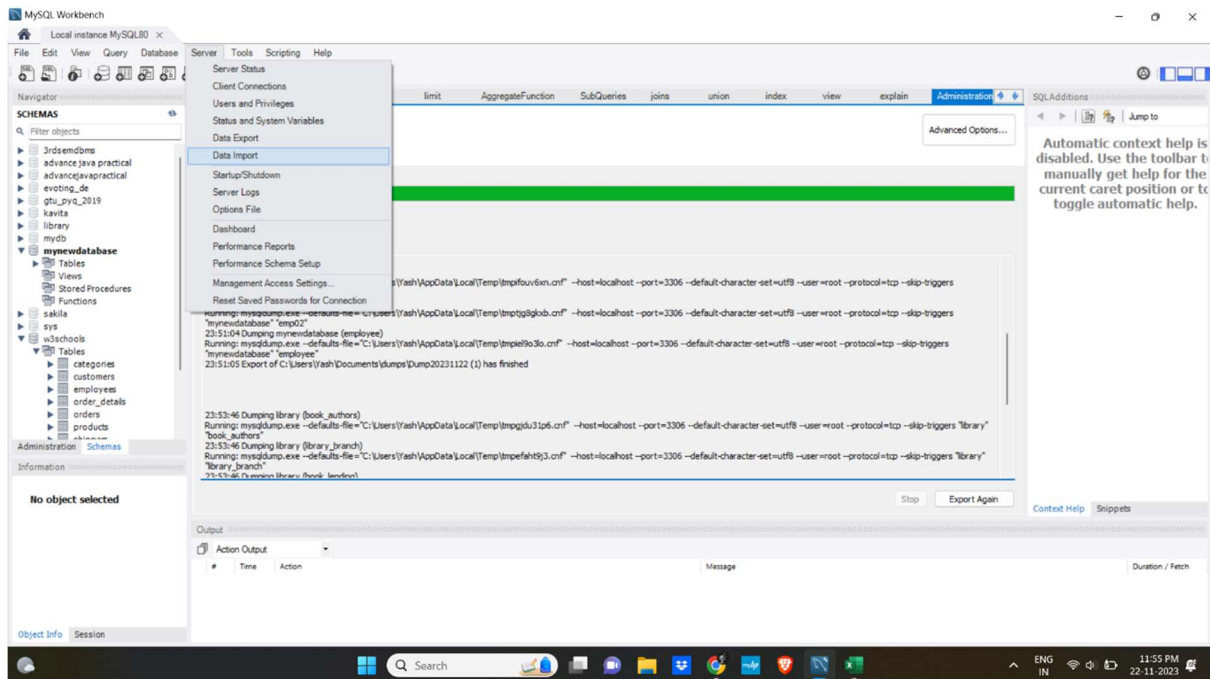
Backing up a MySQL database is crucial for data protection and disaster recovery. The mysqldump utility is commonly used to create backups in MySQL.





Restore

Restoring a MySQL database involves recreating the database from a backup.



Explain

The EXPLAIN keyword is used to obtain information about how the MySQL optimizer executes a SQL query. It provides insights into the execution plan that MySQL has chosen for a specific query, including details about the indexes used, the order of table access, and optimization strategies.

- **id:**
An identifier for the query within the execution plan.
- **select_type:**
The type of SELECT query (e.g., SIMPLE, PRIMARY, SUBQUERY).
- **table:**
The table referenced in the output row.
- **type:**
The type of join that MySQL has chosen for the query (e.g., index scan, full table scan).
- **possible_keys:**
The indexes that MySQL could potentially use.
- **key:**
The index that MySQL has chosen to use.
- **rows:**
The estimated number of rows MySQL expects to examine.

```

EXPLAIN SELECT
    CustomerID,
    CustomerName,
    City
FROM
    Customers
WHERE
    CustomerID > 85;

```

	id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
►	1	SIMPLE	Customers	<small>NULL</small>	range	PRIMARY,Index_Contact,Index_ID	PRIMARY	4	<small>NULL</small>	7	100.00	Using where